

Closed

Opened 1 year ago by  **Ahmad Sherif**

# [RCA] 502s on GitLab.com post Patroni migration maintenance window

## Summary

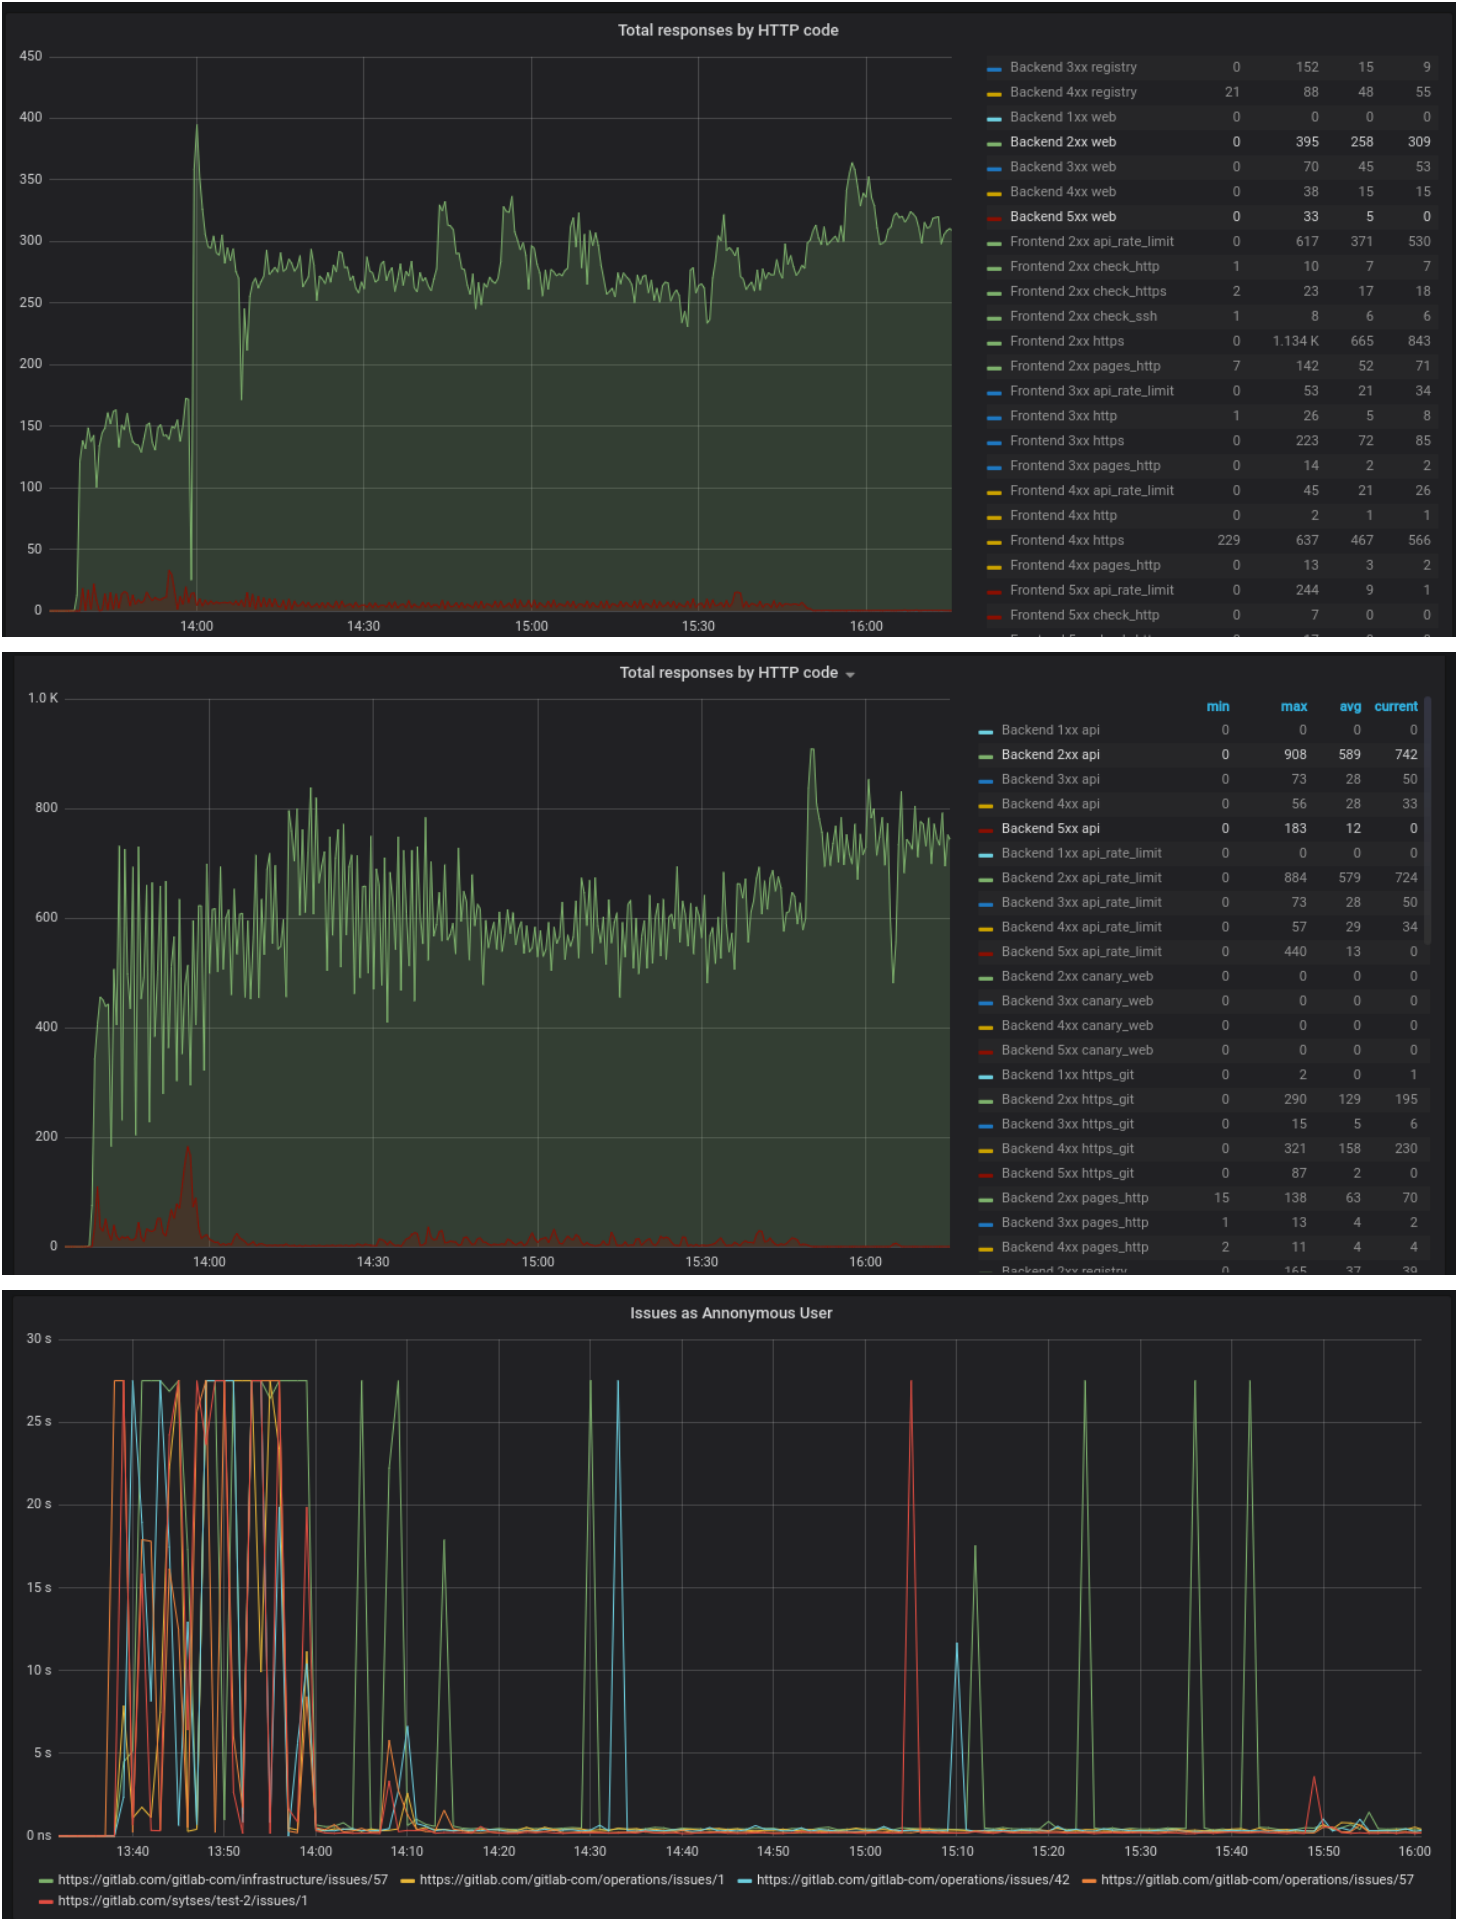
After allowing access to GitLab.com after the Patroni migration, we started noticing slow requests, some ultimately ended with 502 errors.

Service(s) affected : GitLab.com Team attribution : SRE Minutes downtime or degradation : 80 minutes

## Impact & Metrics

Start with the following:

- What was the impact of the incident? (i.e. service outage, sub-service brown-out, exposure of sensitive data, ...)
  - Small percentage of slow requests and 502 errors
- Who was impacted by this incident? (i.e. external customers, internal customers, specific teams, ...)
  - Users of GitLab.com



## Detection & Response

Start with the following:

- How was the incident detected?
  - We saw many Sentry errors like "PG::ConnectionBad: could not connect to server: Connection refused", so we checked pgbouncer logs on the master to see entries like "ERROR accept() failed: Too many open files"
- Did alarming work as expected?
  - No. pgbouncer errors wasn't caught by mtail because they don't match the pattern we match against
- How long did it take from the start of the incident to its detection?
  - Around 5 minutes
- How long did it take from detection to remediation?
  - Initial fix on master took around 15 minutes.
- Were there any issues with the response to the incident? (i.e. bastion host used to access the service was not available, relevant team memeber wasn't page-able, ...)
  - No.

## Root Cause Analysis

- pgbouncer was running with a limit of 1024 (system default) maximum open file descriptors, but was configured to accept a maximum of 2048 client connections
- The Patroni cluster wasn't properly stress-tested before the migration

- with bumped limits

- Identification of the error cause was quick; Sentry errors gave clear hints of where to look

- When entrusting third-parties to do some of sort of testing on an infrastructure component, we should be

- aware of what tests are being made, are these tests enough, and what are the results of these tests

- #5736 (closed).

- ## Blameless RCA Guideline

## Linked issues ?

## Relates to

-  [2018-12-08 Post Maintenance Window 502s on GitLab.com](#)

production#606

- pgbouncer mtail doesn't catch some errors

#5736

🕒 2018-12-30 ... 👤 1 🖼️

- Services running on the Patroni cluster need to have their limits matching omnibus

#5737

🕒 Dev & Ops: 2... 👤 3 🖼️



**Ahmad Sherif** [@ahmadsherif](#) changed the description 1 year ago



Ahmad Sherif [@ahmadsherif](#) · [1 year ago](#)

Maintainer

cc/ [@dawsmith](#) [@Finotto](#)



[Ahmad Sherif @ahmadsherif](#) changed title from **[RCA] Post Maintenance Window 502s on GitLab.com** to **[RCA] 502s on GitLab.com post Patroni migration maintenance window** [1 year ago](#)



**Ahmad Sherif** [@ahmadsherif](#) mentioned in issue [production#606 \(closed\)](#). 1 year ago



**David Smith** 🌴 @dawsmith marked this issue as related to [production#606 \(closed\)](#). 1 year ago



**David Smith** 🌴 @dawsmith marked this issue as related to [#5736 \(closed\)](#) 1 year ago



**David Smith** 🌴 @dawsmith marked this issue as related to [#5737 \(closed\)](#), 1 year ago

Anthony Sandoval @AnthonySandoval · 11 months ago

Owner



[@ahmadsherif](#) I'm going to close this issue. There's an outstanding [corrective action](#) item, but otherwise I don't believe there's anything left to do. Please reopen if you believe otherwise. Thanks.



**Anthony Sandoval** [@AnthonySandoval](#) closed [11 months ago](#)



**ops-gitlab-net** [@ops-gitlab-net](#) mentioned in issue [#7166 \(closed\)](#) [11 months ago](#)

Please [register](#) or [sign in](#) to reply