


Closed

Opened 3 months ago by  **Michal Wasilewski**

Incident review: 2020-03-12 customers.gitlab.com is down

Incident: gitlab-com/gl-infra/production#NNN

Summary

A user ran the customers console as root and this generated a logfile that was owned by root. This log file ownership change cause chef to not run, and the customers app to fail to reply to requests.

- Service(s) affected : customers.gitlab.com
- Team attribution :
- Minutes downtime or degradation : 15 minutes

For calculating duration of event, use the [Platform Metrics Dashboard](#) to look at appdex and SLO violations.

Impact & Metrics

Start with the following:

- What was the impact of the incident? (i.e. service outage, sub-service brown-out, exposure of sensitive data, ...)
- Who was impacted by this incident? (i.e. external customers, internal customers, specific teams, ...)
- How did the incident impact customers? (i.e. preventing them from doing X, incorrect display of Y, ...)
- How many attempts were made to access the impacted service/feature?
- How many customers were affected?
- How many customers tried to access the impacted service/feature?

Include any additional metrics that are of relevance.

Provide any relevant graphs that could help understand the impact of the incident and its dynamics.

Detection & Response

Start with the following:

- How was the incident detected?
- Did alarming work as expected?
- How long did it take from the start of the incident to its detection?
- How long did it take from detection to remediation?
- Were there any issues with the response to the incident? (i.e. bastion host used to access the service was not available, relevant team memeber wasn't page-able, ...)

Root Cause Analysis

1. The log file was read-only to the customers application.
2. The log file had been written to, or created by, a root owned console application.
3. A user had followed instructions that specified to use `root` , and not the `gitlab-customers` user.
4. The documentation was old/wrong.

What went well

Start with the following:

- Identify the things that worked well or as expected.
- Any additional call-outs for what went particularly well.

What can be improved

Start with the following:

- Using the root cause analysis, explain what can be improved to prevent this from happening again.
- Is there anything that could have been done to improve the detection or time to detection?
- Is there anything that could have been done to improve the response or time to response?

- Is there an existing issue that would have either prevented this incident or reduced the impact?
- Did we have any indication or beforehand knowledge that this incident might take place?

Corrective actions

- List issues that have been created as corrective actions from this incident.
- For each issue, include the following:
 - - Issue labeled as `corrective action` .
 - Include an estimated date of completion of the corrective action.
 - Incldue the named individual who owns the delivery of the corrective action.
- [#9498 \(closed\)](#).

Guidelines

- [Blameless RCA Guideline](#)
- [5 whys](#)

Edited 3 months ago by [Cameron McFarland](#)

Linked issues ⓘ

2

Relates to

[2020-03-12 customers.gitlab.com is down](#)

production#1763

[Update customers.gitlab.com cheat sheet to prevent outages](#)

#9498

Reliability - M...

1

[Michal Wasilewski](#) [@mwasilewski-gitlab](#) added [IncidentReview](#) label [3 months ago](#)

[Michal Wasilewski](#) [@mwasilewski-gitlab](#) · [3 months ago](#)

Maintainer

[@cmcfarland](#) Would you mind taking a look at what potentially changed permissions on that file so that it doesn't happen again? This MR looks related to permissions used by logrotate: https://gitlab.com/gitlab-cookbooks/cookbook-customers-gitlab-com/-/merge_requests/80/diffs . Maybe it's something that happens as part of uploading files to GCS?

Collapse replies

[Cameron McFarland](#) [@cmcfarland](#) · [3 months ago](#)

Maintainer

Yes. I'll look into this.

[Cameron McFarland](#) [@cmcfarland](#) · [3 months ago](#)

Maintainer

It looks like most of the correction for this is taking place in this thread: [production#1763 \(comment 303875051\)](#).

Please [register](#) or [sign in](#) to reply

[Michal Wasilewski](#) [@mwasilewski-gitlab](#) mentioned in issue [on-call-handovers#415 \(closed\)](#), [3 months ago](#)

[Cameron McFarland](#) [@cmcfarland](#) assigned to [@cmcfarland](#) [3 months ago](#)

[Cameron McFarland](#) [@cmcfarland](#) marked this issue as related to [production#1763 \(closed\)](#), [3 months ago](#)

[Cameron McFarland](#) [@cmcfarland](#) changed the description [3 months ago](#) · [Compare with previous version](#) ▾

[Cameron McFarland](#) [@cmcfarland](#) · [3 months ago](#)

Maintainer

I think we should update the [documentation](#) that [@arihantar](#) referenced to help make sure this doesn't happen again. The [customers readme](#) does have the right information.



Cameron McFarland [@cmcfarland](#) · 3 months ago

Maintainer

I think more actions beyond this one are not worth the effort. Monitoring and logging for customers could be better, but with it being a node in azure and our larger goal of getting it into auto-devops, effort should go into the migration before more improvements.



Cameron McFarland [@cmcfarland](#) mentioned in issue [#9498 \(closed\)](#), 3 months ago



Cameron McFarland [@cmcfarland](#) marked this issue as related to [#9498 \(closed\)](#), 3 months ago



Cameron McFarland [@cmcfarland](#) changed the description 3 months ago · [Compare with previous version](#) ▾



Cameron McFarland [@cmcfarland](#) closed 3 months ago



ops-gitlab-net [@ops-gitlab-net](#) mentioned in issue [#9472 \(closed\)](#), 3 months ago



Devin Sylva [@devin](#) · 3 months ago

Maintainer

This problem happened again just now. Fixed the same way.

Please [register](#) or [sign in](#) to reply