


Closed

Opened 1 month ago by  Igor

nonprod-log cluster unavailable

Our nonprod-log cluster is currently unavailable.

It stopped ingesting around Friday 2020-05-23 13:30:



The monitoring cluster says everything is fine, though the ES cloud ui does not:

Unhealthy deployment

×

An Elasticsearch shard unavailable

🔍

Last successful snapshot was 2 days ago

Trying to access the cluster at <https://nonprod-log.gitlab.net/> gives us the "unhealthy backend" error from google cloud's load balancer:

Error: Server Error

The server encountered a temporary error and could not complete your request.

Please try again in 30 seconds.

Error: Server Error

The server encountered a temporary error and could not complete your request.

Please try again in 30 seconds.

This points to either the IAP or an unhealthy endpoint behind it.

Trying to access the cluster directly at <https://e8eb2b813ecf40e3ab39829e7cd3a4d9.us-central1.gcp.cloud.es.io:9243/> does load the login page, but upon login, we get:

```
{"statusCode":503,"error":"Service Unavailable","message":"[unavailable_shards_exception] a
```

I believe what has happened here:

- `.security-7` is a system index that holds user information
- We lost that index, and now logins are broken everywhere

Potential quickfix that I will try:

- Restore that one index from a snapshot.

Longer-term fix:





- Ensure the system indices have at least one replica shard configured, so that losing a single shard does not mean a complete outage.
- File upstream issue about monitoring cluster not warning about this.


cc [@mwasilewski-gitlab](#)

Edited 1 month ago by [Igor](#)

Linked issues ?

0

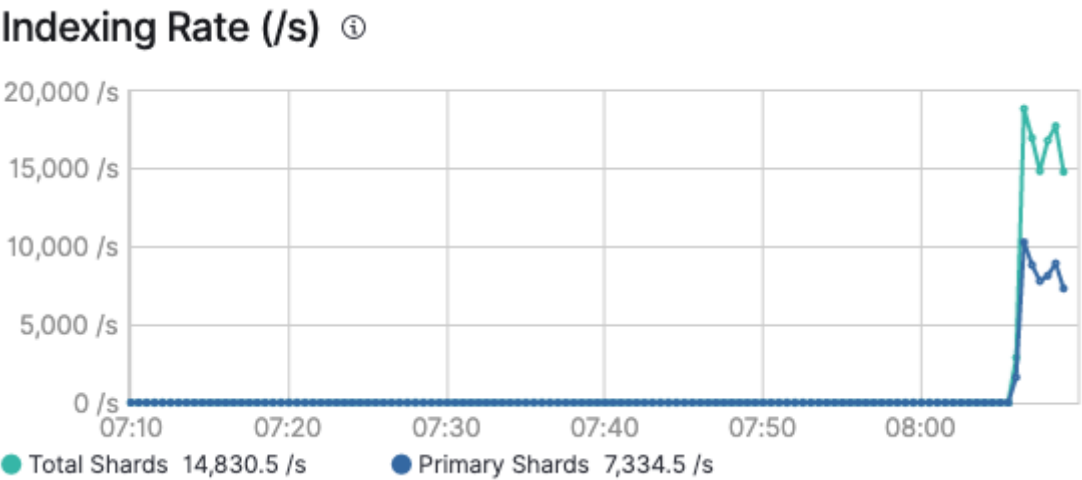
- [Igor @igorwww](#) changed milestone to [%Observability Team - May 2020](#) 1 month ago
- [Igor @igorwww](#) added [team Observability](#) [Service ELK](#) scoped labels 1 month ago
- [Igor @igorwww](#) added [toil](#) label 1 month ago
- [Igor @igorwww](#) changed the description 1 month ago · [Compare with previous version](#)


[Igor @igorwww](#) · 1 month ago Owner

I deleted the existing unhealthy `.security-7` index and restored it from snapshot.

Kibana is now accessible again.

Indexing also started again (at 18k/s, normal rate is 1k/s):




[Igor @igorwww](#) · 1 month ago Owner

It looks like we ~~do~~ have a replica shard configured for `.security-7`. So it looks like we may have lost both the primary and the replica in this case.


Since it's a critical index for everything (and it's tiny), perhaps it makes sense to increase the replication factor.


▼ Collapse replies

[Michal Wasilewski @mwasilewski-gitlab](#) · 1 month ago Maintainer

There's only 2 hot and 2 warm nodes in the cluster, so if we wanted to have more than 2 shards for the index we would have to resize the cluster.

Please [register](#) or [sign in](#) to reply

[Igor @igorwww](#) added [workflow-infra](#) [In Progress](#) scoped label 1 month ago

[Igor @igorwww](#) · 1 month ago Owner

Closing this, since the cluster has been recovered. If it happens more frequently we can look at further measures.

[Igor @igorwww](#) closed 1 month ago




Igor @igorwwwwwwwwwwwwwwwwwwww added workflow-infra Done scoped label and automatically removed workflow-infra In Progress label 1 month ago



Igor @igorwwwwwwwwwwwwwwwwwwww · 1 month ago

Owner

And just to close the loop here, the backlog was replayed successfully:



Please [register](#) or [sign in](#) to reply