



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#).
Please help us clean up older documentation referring to tools.wmflabs.org!

Incident documentation/20170512-API

[< Incident documentation](#)

Contents [\[hide\]](#)

- [1 Summary](#)
- [2 Timeline](#)
- [3 Other service impacts because of the core change](#)
- [4 Conclusions](#)
- [5 Actionables](#)

Summary

The TextExtracts extension (which provides the `extracts` prop to the MediaWiki API) was providing empty `extract` fields for queries for a period of 18 hours. This affected RESTBase and subsequently mobileapps.

Timeline

- 14:26:30:** Icinga begins reporting problems with RESTBase

```
PROBLEM - restbase endpoints health on restbase1007 is CRITICAL:
/en.wikipedia.org/v1/page/random/{format} (Random title redirect) is
CRITICAL: Test Random title redirect returned the unexpected status 503
(expecting: 303): /en.wikipedia.org/v1/page/summary/{title} (Get summary from
storage) is WARNING: Test Get summary from storage responds with unexpected
body: /extract = : /en.wikipedia.org/v1/page/title/{title} (Get rev by title
```

- 14:27:34:** `bblack`, `_Joe_`, `godog`, `elukey`, and `mobrovac` begin investigation in IRC
- 15:06:40:**

```
<mobrovac> the "extract" field is empty for any page i try to fetch the summary
for
<mobrovac> we get that from the MW API
```

- 15:08:14:** `thcipriani` is pinged to rollback from `1.30.0-wmf.1` → `1.29.0-wmf.21` since API has been pinpointed, `thcipriani` rolls back all wikis on `mwdebug1002` to `1.29.0-wmf.21` and pings `mobrovac` and `_Joe_` to test.
- 15:24:19:** [RESTBase query](#) is tried and determined that a rollback does not fix it, but rollback + null edit does, caching is involved
- ~15:27:** Help is needed from Mediawiki devs, `ebernhardson` troubleshoots `anomie` is in meeting but helps, too
- 15:47:58:**

```
<ebernhardson> hmm, so i can verify at this point that
ExtractFormatter::filterContent() is the part that's stripping everything, but
not sure yet what exactly
<ebernhardson> a full page goes in, and what comes out is just a body tag and an
html comment
```

- 15:58:00:** problem found

[Main page](#)
[Recent changes](#)
[Server admin log \(Prod\)](#)
[Server admin log \(RelEng\)](#)
[Deployments](#)
[SRE/Operations Help](#)
[Incident status](#)

[Cloud VPS & Toolforge](#)
[Cloud VPS documentation](#)
[Toolforge documentation](#)
[Request Cloud VPS project](#)
[Server admin log \(Cloud VPS\)](#)

[Tools](#)
[What links here](#)
[Related changes](#)
[Special pages](#)
[Permanent link](#)
[Page information](#)
[Cite this page](#)

[Print/export](#)
[Create a book](#)
[Download as PDF](#)
[Printable version](#)

```
<ebernhardson> so, after manually removing the <div class="mw-parser-output">
from the beginning, we now get the correct response from
https://en.wikipedia.org/w/api.php?
action=query&format=json&prop=extracts&formatversion=2&redirects=true&titles=Kuiper
belt (because it's cached in memcached from my mwrepl session)
```

```
16:13:08:    <ebernhardson> so random friday options: #1) div.mw-paser-output ->
some other tag? not sure what though #2) revert div.mw-parser-output #3) Drop
'div' from list of things to strip from text extracts, live with some extra
content that wasn't wanted
16:15:05:    <thcipriani> option 2 is the one that seems like not new code :)
16:32:19:    <mobrovac> anomie, ebernhardson, thcipriani: at this point, I
honestly prefer #2, it seems the safest, the rest can be dealt with during
quieter and more appropriate times
16:33:03:    * anomie submitted https://gerrit.wikimedia.org/r/#/c/353565/ and
https://gerrit.wikimedia.org/r/#/c/353566/ FWIW
16:41:09:    <thcipriani> anomie: so the options to fix would be revert
https://gerrit.wikimedia.org/r/#/c/350634/ or merge these 2
https://gerrit.wikimedia.org/r/#/q/topic:bug/T165161 ?
16:42:37:    <anomie> thcipriani: Yes. Either of those options should also fix
T165115. In both cases additional caching being done by
TextExtracts/MobileFrontend might mask the fix.
16:52:23:    <thcipriani> mobrovac: anomie ok, I'm going to revert
https://gerrit.wikimedia.org/r/#/c/350634/ and run a sync since (a) it's friday
before a week where releng won't be looking deployment stuffs and I don't want
to deploy new code and (b) it looks like it was supposed to be reverted for
wmf.1 anyway https://gerrit.wikimedia.org/r/#/c/352577/ Is there any concern
about a revert breaking anything additional
17:01:58:    <anomie> thcipriani: I don't think reverting in 1.30.0-wmf.1 will
break anything else.
```

- **18:10:** thcipriani@tin: Finished scap: Revert "Wrap parser output in" 4/4 (duration: 19m 13s)
- **18:20:** legoktm reports revert not working submits <https://gerrit.wikimedia.org/r/#/c/353593/> to version the memcached key so that the cache will be cleared
- **18:27:** code live on mwdebug1002, but determined not to be working, MaxSem explains: **18:35:37:** <MaxSem> however, due to reverts of the whole wrapping diff thing, we have corrupt cache
- **18:43:** legoktm creates <https://gerrit.wikimedia.org/r/#/c/353596/> to clear parser cache and it's deployed, scap catches problem and <https://gerrit.wikimedia.org/r/#/c/353597/> is created and sync to correct problem
- **19:17:** thcipriani syncs <https://gerrit.wikimedia.org/r/#/c/353593/> to change memcached key

Other service impacts because of the core change

- **T165139** - inline <math> editing broken in VE
- **T165070** - spurious linter errors reported for <maplink> tags

Underlying cause for both of these is the extra <div> tag around Parsoid's extension output because Parsoid calls the mediawiki API (action=parse) to process extension tags. This has since been addressed in <https://gerrit.wikimedia.org/r/#/c/353707/> and <https://gerrit.wikimedia.org/r/#/c/354050/> so that when the core change is redeployed, Parsoid and downstream clients aren't impacted.

Conclusions

- TextExtracts enables a lot of services on which we depend
- TextExtracts can be broken by adding a div to parser output
- There are, evidently, few alarms that catch empty textextract summaries in the API
- This outcome may have been expected judging from a revert that was proposed before the train <https://gerrit.wikimedia.org/r/#/c/352577/>
- The core patch was proposed to be reverted in master as part of <https://gerrit.wikimedia.org/r/#/c/352577/> but an alternative proposal was to not revert it in master, but to revert it in 1.30.0-wmf.1 instead. But, it appears that the revert didn't happen in wmf.1 before deploy.
- Parsoid uses the mediawiki API to process extension tags and there a number of downstream services that depend on Parsoid: VE, Flow, Content Translation, Android App (via Mobile Content Service), and Linter. This bug would have showed up in the beta cluster as reported in T165139. So, there is clearly a testing gap here.

Actionables

- Better alarm mechanism for empty textextract summaries so they don't languish for 18 hours
- (/me is likely missing lots of subtlety, but armchair-quarterbacks anyway) Adding a div shouldn't break textextract, is there some kind of class/id (e.g., `.mw-ext-textextract-start`) that can be added to parser output that can then be used by textextract that will have the single express purpose of being used by textextract? Adding a div breaking stuff makes the mechanisms by which it is extracting data seem fragile.
- If it is likely that deploying a change will cause breakage (i.e., <https://gerrit.wikimedia.org/r/#/c/352577/>) let deployers/branchers know so that they can backport fixes or reverts before problems manifest.
- Identify how to update the QA tests that are run on the beta cluster so that any potential breakage of VE, etc. is caught before changes are rolled out to production.

Category: [Incident documentation](#)

This page was last edited on 16 May 2017, at 21:27.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.

[Privacy policy](#) [About](#)

[Disclaimers](#) [Code of Conduct](#) [Developers](#) [Statistics](#) [Cookie statement](#) [Mobile view](#)

[Wikitech](#)

