



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#).  
Please help us clean up older documentation referring to tools.wmflabs.org!

# Incident documentation/20170711-EventLogging

[< Incident documentation](#)

## Contents [\[hide\]](#)

- 1 [Summary](#)
- 2 [Timeline](#)
  - 2.1 [Mon July 10 11:31:10 UTC 2017](#)
  - 2.2 [Mon Apr 6 18:32 UTC 2015](#)
- 3 [Conclusions](#)
- 4 [Actionables](#)

## Summary

Within the previous month, we started importing EventBus generated events into the MySQL EventLogging Analytics database(s) for <https://phabricator.wikimedia.org/T150369>. These events have a different schema format and schema repository than the original EventLogging Analytics events.

EventBus style events are versioned just like the EventLogging Analytics ones, but thus far, we had not bumped any version numbers for backwards compatible schema changes. However, now that MySQL is being used for these events, and MySQL tables are not automatically evolved, we must bump EventBus style schemas from now on. This was not done for a change that added a `rev_content_changed` field (<https://gerrit.wikimedia.org/r/#/c/362321/>). When events with this new field started flowing in from EventBus, the eventlogging process that was attempting to insert these events died. It was immediately restarted by upstart, but it flapped every 5 minutes, every time a batch of `mediawiki_page_create` events were attempted to be inserted.

This caused minimal data loss from other unrelated schemas. These events were then backfilled from secondary log files into MySQL. Along the way, Otto forgot that we were now filtering out bot events, and inserted over 2 million extra events into the MySQL databases for the days of 2017-07-10 - 2017-07-12.

## Timeline

### Mon July 10 11:31:10 UTC 2017

<https://gerrit.wikimedia.org/r/#/c/362322/> is merged. This is slowly deployed to wikis as part of the regular deploy train.

### Mon Apr 6 18:32 UTC 2015

Roan files <https://phabricator.wikimedia.org/T170486>. Otto investigates, realizes what happening. He then alters the existing eventbus event tables in MySQL to add the missing `rev_content_changed` field to allow insertions to happen. He also merges a puppet change (<https://gerrit.wikimedia.org/r/#/c/364881/>) to separate the processes responsible for inserting EventLogging Analytics events, and events from EventBus to prevent something like this from happening again.

## Conclusions

EventBus event schemas from here on out need to have version numbers bumped when we modify them.

## Actionables

- Status: ■ **Unresolved** Monitor if eventlogging processes flap

Categories: [Incident documentation](#) | [EventLogging/Incident documentation](#)

[Main page](#)  
[Recent changes](#)  
[Server admin log \(Prod\)](#)  
[Server admin log \(RelEng\)](#)  
[Deployments](#)  
[SRE/Operations Help](#)  
[Incident status](#)

[Cloud VPS & Toolforge](#)  
[Cloud VPS documentation](#)  
[Toolforge documentation](#)  
[Request Cloud VPS project](#)  
[Server admin log \(Cloud VPS\)](#)

[Tools](#)  
[What links here](#)  
[Related changes](#)  
[Special pages](#)  
[Permanent link](#)  
[Page information](#)  
[Cite this page](#)

[Print/export](#)  
[Create a book](#)  
[Download as PDF](#)  
[Printable version](#)

This page was last edited on 17 July 2017, at 19:12.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.

[Privacy policy](#) [About](#)

[Disclaimers](#) [Code of Conduct](#) [Developers](#) [Statistics](#) [Cookie statement](#) [Mobile view](#)

[Wikitech](#)

