



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#) .  
Please help us clean up older documentation referring to tools.wmflabs.org!

# Incident documentation/20170920-Elasticsearch

[< Incident documentation](#)

## Contents [\[hide\]](#)

- [1 Summary](#)
- [2 Timeline](#)
- [3 Conclusions](#)
- [4 Actionables](#)

## Summary

Yesterday, we did a cold restart of the elasticsearch / cirrus eqiad cluster. This restart did not go as planned. It did not generate any user facing impact, since we moved all the traffic to codfw before the restart. It did impact logstash (more of that in a different report).

## Timeline

- 13:15: search traffic moved to codfw
- 15:36: start of plugin upgrade and cluster restart on elasticsearch eqiad
- 15:46: shutdown and restart of all elasticsearch nodes, cluster is not recovering
- investigation (detailed before, no exact timeline)
- 16:34: restart of elastic1051, cluster is starting to recover, cluster is in yellow state after a few minutes (ready to serve traffic)
- 21:00 (approx): full recovery of cluster

## Conclusions

elastic1051 did not restart cleanly, due to failed cleanup of previous plugins. This was not investigated right away, as we thought that recovering the cluster was more urgent and that the cluster should recover without issue even if one node is unavailable. We now know differently.

The parameters governing the recovery of the cluster are:

- recover\_after\_nodes: 24
- recover\_after\_time: 20m
- expected\_nodes: 36

Our previous understanding was that the cluster would start recovering with 24 of the 36 nodes available, or start recovering anyway after 20 minutes. So we were not expecting that one node being down would prevent a recovery. Our current understanding is that elasticsearch will wait for `expected_nodes` to be available, and recover anyway after `recover_after_time` if at least `recover_after_nodes` are present.

## Actionables

- Update [our documentation](#).
- Tune recovery settings [Task T176409](#)

Category: [Incident documentation](#)

This page was last edited on 21 September 2017, at 12:59.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.

[Main page](#)  
[Recent changes](#)  
[Server admin log \(Prod\)](#)  
[Server admin log \(RelEng\)](#)  
[Deployments](#)  
[SRE/Operations Help](#)  
[Incident status](#)

[Cloud VPS & Toolforge](#)

[Cloud VPS documentation](#)

[Toolforge documentation](#)

[Request Cloud VPS project](#)

[Server admin log \(Cloud VPS\)](#)

[Tools](#)

[What links here](#)

[Related changes](#)

[Special pages](#)

[Permanent link](#)

[Page information](#)

[Cite this page](#)

[Print/export](#)

[Create a book](#)

[Download as PDF](#)

[Printable version](#)

