



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#).
Please help us clean up older documentation referring to tools.wmflabs.org!

Incident documentation/20180214-labvirt1008-failure

[< Incident documentation](#)

Contents [\[hide\]](#)

- [1 Summary](#)
- [2 Timeline](#)
- [3 Conclusions](#)
- [4 Actionables](#)

Summary

Labvirt1008 seems to have overheated and gone down. This effects tenants as well as virtual VPS infrastructure

Timeline

- 7:20 UTC labvirt1008 rebooted
- 8:00 UTC moritzm logged a task about it
- There was no paging or alerting. This is a problem.
- 11:00 UTC Chase woke up and started investigating the extent of the outage, and looking for impact on Toolforge especially
- 12:58 UTC Chase sent an email about impact to cloud-announce <https://lists.wikimedia.org/pipermail/cloud-announce/2018-February/000023.html> with a list of [affected instances](#).

Conclusions

We know that our instance storage is local and ephemeral. We should ensure that is documented for tenants in easy to find places, and re-ensure that our mechanism that keep critical redundant components spread across labvirt's are working. In our world though a single hypervisor is a special snowflake and I believe we should have been paged on this outage, but seem not to have been. It was my understanding that a full instance storage partition should have paged if nothing else, and in this case the failure of that check.

Actionables

- Coordinate with DC OPS to deal with overheating [phab:T187292](#)
- Look at moving tenant instances to another labvirt (we should have a standing spare)
- Investigate what should have paged and why it did not (and fix it)

Category: [Incident documentation](#)

This page was last edited on 19 October 2018, at 11:50.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.