



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#).  
Please help us clean up older documentation referring to tools.wmflabs.org!

# Incident documentation/20180226-WikibaseQualityConstraints

[< Incident documentation](#)

## Contents [\[hide\]](#)

- [Summary](#)
- [Timeline](#)
- [Conclusions](#)
- [Actionables](#)

## Summary

When caching of constraint check results was enabled on Wikidata, some API requests would trigger SQL queries on the `page`, `revision` and `text` tables without any `WHERE` clause. This resulted in extremely high server load of general-purpose replica servers in s8.

## Timeline

All times in UTC.

- 14:22 – zeljkof deploys [413724](#) to mwdebug1002 as part of SWAT.
- 14:22-14:26 – Lucas Werkmeister ([Lucas\\_WMDE](#)) tests the change on Wikidata. It seems to work as expected – successive calls of `wbcheckconstraints` on the same entity are much faster, since the result is read from cache.
- 14:27 – Lucas Werkmeister ([Lucas\\_WMDE](#)) reports back to zeljkof and confirms the change is ready to deploy.
- 14:28 – zeljkof deploys the change. [Server admin log entries during this time](#)
- 14:28 – eth0 traffic on db1109 immediately begins to grow, crossing 600 Mbps within the next ten minutes (the normal level is around 50 Mbps) – see [Grafana](#). db1092 and db1104 are also affected. CPU load, number of running processes, and disk I/O also increase accordingly. However, according to jynus, "the slowdown was slow to buildup, so it was not detected by monitoring immediately".
- 17:24 – marostegui and jynus comment on [phabricator:T184812](#), pointing to the massive spike in Wikidata replicas ([Grafana](#)).
- 17:34 – jynus "deploying new query killer to db1109". [Server admin log entries during this time](#)
- 17:41 – Chad deploys a revert of the config change ([414720](#)). [log entry](#) (stashbot failed to log this and some other messages around that time to wiki)
- 17:41 – eth0 traffic, CPU load etc. begin to drop again ([Grafana](#)).
- 17:41 – jynus "I am going to kill [queries?] [a]gain to see they don't come back"
- 17:44 – eth0 traffic, CPU load etc. back to normal levels ([Grafana](#)).
- 17:45 – jynus "things seem under control"

## Conclusions

Due to a logic bug ([phabricator:T188384](#)), WikibaseQualityConstraints' `CachingResultsBuilder` asked Wikibase' `WikiPageEntityMetaDataTableLookup` for the latest revision information of an empty list of entities. The `WikiPageEntityMetaDataTableLookup` had a special safeguard (added in [294340](#)) for this case to avoid costly queries, adding a condition `"0"`. However, `Database::selectSQLText` actually turns this condition into a query with no `WHERE` clause (instead of something like `WHERE FALSE`, which `WikiPageEntityMetaDataTableLookup` probably intended). WikibaseQualityConstraints should pay more attention to special cases when requesting entity IDs (empty list, long list); Wikibase should make sure that safeguards actually work as intended; and core should not let simple programming errors result in completely unlimited queries.

[Main page](#)  
[Recent changes](#)  
[Server admin log \(Prod\)](#)  
[Server admin log \(RelEng\)](#)  
[Deployments](#)  
[SRE/Operations Help](#)  
[Incident status](#)

[Cloud VPS & Toolforge](#)

[Cloud VPS documentation](#)

[Toolforge documentation](#)

[Request Cloud VPS project](#)

[Server admin log \(Cloud VPS\)](#)

Tools

[What links here](#)

[Related changes](#)

[Special pages](#)

[Permanent link](#)

[Page information](#)

[Cite this page](#)

Print/export

[Create a book](#)

[Download as PDF](#)

[Printable version](#)

It is also problematic that the incident apparently went undetected for almost three hours, even though it was visible in Grafana within minutes of deployment, but I don't know why this was possible, or what should have prevented it.

## Actionables

- WikibaseQualityConstraints: only filter for result statuses after collecting metadata. [phab:T188384](#)
- WikibaseQualityConstraints: don't pass an empty list of entity IDs to the `WikiPageEntityMetadataAccessor`. [phab:T188311](#)
- WikibaseQualityConstraints: don't pass overly *long* lists of entity IDs there, either. [phab:T188312](#)
- WikibaseQualityConstraints + Wikibase: don't join `page`, `revision` and `text` when we just need `page.page_latest`. [phab:T188381](#)
- Wikibase: fix the safeguard for an empty list of conditions in `WikiPageEntityMetadataLookup`. [phab:T188313](#)
- core: only `$conds = []` should result in a missing `WHERE` clause, not any other value that is `empty` according to the PHP `empty` function. [phab:T188314](#)
- monitoring: detect high server load earlier – jynus "a prometheus alert would be nice" [phab:T188317](#)

Category: [Incident documentation](#)

This page was last edited on 19 October 2018, at 11:53.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.

[Privacy policy](#) [About](#)

[Disclaimers](#) [Code of Conduct](#) [Developers](#) [Statistics](#) [Cookie statement](#) [Mobile view](#)

[Wikitech](#)

