



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#).  
Please help us clean up older documentation referring to tools.wmflabs.org!

# Incident documentation/20181001-DeleteLocalPasswords

[< Incident documentation](#)

## Contents [\[hide\]](#)

- [1 Summary](#)
- [2 Timeline](#)
- [3 Conclusions](#)
- [4 Links to relevant documentation](#)
- [5 Actionables](#)

## Summary

The MediaWiki maintenance script ([T201009](#) / [T201009](#), [DeleteLocalPasswords.php](#)) did not include a `waitForReplication` call. When run on certain databases (commons and dewiki), it updated over a million rows, causing an extended period read only time and slowdown of requests on Commons (and other projects using Commons database resources) approximately between 10:48 and 11:01 and on the German Wikipedia between 11:02 and 11:16 UTC.

## Timeline

[SAL](#) for 2018-10-01: Script starts (this wasn't known at the time of the initial response)

```
10:10 <Amirl> mwscript
extensions/CentralAuth/maintenance/deleteLocalPasswords.php --wiki=fawiki --
delete (T201009)
10:15 <Amirl> ladsgroup@mwmaint2001:~$ mwscript
extensions/CentralAuth/maintenance/deleteLocalPasswords.php --prefix on all
CentralAuth wikis (T201009)
```

First alert on IRC:

```
10:51:16 <icinga-wm> PROBLEM - MariaDB Slave Lag: s4 on db2058 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 483.80 seconds
```

(thinking it is a server-only issue, and seeing no slow query ongoing but 'slow' loadgroup queries db2058 is depooled and the 2 slow queries killed)

```
10:58:57 <logmsgbot> !log jynus@deploy1001 Synchronized wmf-config/db-
codfw.php: Depool db2058 (duration: 00m 57s)
```

But other hosts starts alerting, too:

```
11:04:46 <icinga-wm> PROBLEM - MariaDB Slave Lag: s4 on db2090 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 310.07 seconds
11:05:37 <icinga-wm> PROBLEM - MariaDB Slave Lag: s4 on db2073 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 301.54 seconds
11:05:56 <icinga-wm> PROBLEM - MariaDB Slave Lag: s4 on db2084 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 318.34 seconds
```

Swat is paused, probably a master or code bug issue is thought. Main queries happening at the moment from replication are DeleteLocalPassword ones from mwmaint2001. Discussion happens if to kill the commons script,

[Main page](#)[Recent changes](#)[Server admin log \(Prod\)](#)[Server admin log \(RelEng\)](#)[Deployments](#)[SRE/Operations Help](#)[Incident status](#)[Cloud VPS & Toolforge](#)[Cloud VPS documentation](#)[Toolforge documentation](#)[Request Cloud VPS project](#)[Server admin log \(Cloud VPS\)](#)[Tools](#)[What links here](#)[Related changes](#)[Special pages](#)[Permanent link](#)[Page information](#)[Cite this page](#)[Print/export](#)[Create a book](#)[Download as PDF](#)[Printable version](#)

and it is done (but the issue just jumps into dewiki).

```
11:06:57 <icinga-wm> RECOVERY - MariaDB Slave Lag: s4 on db2073 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
11:07:16 <icinga-wm> RECOVERY - MariaDB Slave Lag: s4 on db2084 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
11:08:12 <icinga-wm> RECOVERY - MariaDB Slave Lag: s4 on db1097 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
11:08:36 <icinga-wm> RECOVERY - MariaDB Slave Lag: s4 on db2090 is OK: OK
slave_sql_lag Replication lag: 0.05 seconds
11:09:17 <icinga-wm> RECOVERY - MariaDB Slave Lag: s4 on db2058 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
```

Finally, the parent process is found:

```
11:09 <_joe_> killed bash runner.sh by user ladsgroup on mwmain2001
```

dewiki ongoing process is killed at the same time, causing temporary lag, too (but not enough to page):

```
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1102 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 440.58 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1110 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 441.13 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1124 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 445.59 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1082 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 452.27 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1097 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 464.86 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1100 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 466.61 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1096 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 474.90 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1070 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 334.97 seconds
11:17:01 <icinga-wm> PROBLEM - MariaDB Slave Lag: s5 on db1113 is CRITICAL:
CRITICAL slave_sql_lag Replication lag: 490.54 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1110 is OK: OK
slave_sql_lag Replication lag: 59.86 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1124 is OK: OK
slave_sql_lag Replication lag: 0.43 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1082 is OK: OK
slave_sql_lag Replication lag: 0.03 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1097 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1100 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1096 is OK: OK
slave_sql_lag Replication lag: 0.16 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1070 is OK: OK
slave_sql_lag Replication lag: 0.10 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1113 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
11:17:01 <icinga-wm> RECOVERY - MariaDB Slave Lag: s5 on db1102 is OK: OK
slave_sql_lag Replication lag: 0.00 seconds
```

## Conclusions

- Do not forget to add `waitForReplication` calls.
- Registering the maintenance on the [Deployments](#) page might have helped to identify the issue more quickly.
- The script ran successfully on fawiki (medium-sized). Is there some way to notice that a script is causing replag (assuming it did so at all) when it is not large enough to bring things down?

## Links to relevant documentation

Not sure what would be relevant here. [waitForReplication\(\)](#)?

## Actionables

- Consider adding code to the MediaWiki DB layer to detect huge numbers of writes without waiting for slaves - [phab:T205893](#)
- Consider adding dry-run options on all maintenance scripts, even those that seem trivial in the future
- Communicate more clearly to sysadmins (people that can respond quickly to incidents) of what, where and when maintenance scripts are done as documented at <https://wikitech.wikimedia.org/wiki/Deployments> [↗](#) (Long running changes/scripts) so the script can be easily and faster killed if an unexpected issue arises

Category: [Incident documentation](#)

This page was last edited on 19 October 2018, at 12:16.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.

[Privacy policy](#) [About](#)

[Disclaimers](#) [Code of Conduct](#) [Developers](#) [Statistics](#) [Cookie statement](#) [Mobile view](#)

[Wikitech](#)

