Main page
Recent changes
Server admin log (Prod)
Server admin log (RelEng)
Deployments
SRE/Operations Help
Incident status

Cloud VPS & Toolforge
Cloud VPS documentation
Toolforge documentation
Request Cloud VPS project
Server admin log (Cloud VPS)

Tools
What links here
Related changes
Special pages
Permanent link
Page information
Cite this page

Print/export
Create a book
Download as PDF
Printable version

Incident documentation/20190208-logstash-mediawiki

< Incident documentation

Contents [hide]

Summary

A single database host down caused an exceptional influx of log entries from MediaWiki, which in turn caused overload in logstash ingestion.

Timeline

- 12:18 db1114 DOWN alert fires, see also bug T214720
- 12:30 DBAs are engaged
- 12:42 UDP packet loss alerts fire for logstash
- 12:44 db1114 is depooled
- 13:09 UDP packet loss for logstash recovers

Conclusions

MediaWiki alone was able to cause a logstash overload, resulting in UDP packet loss. Applications using UDP as log transport have experienced loss of logs, while applications using the new logging pipeline (i.e. writing to Kafka, MediaWiki included) experienced a slowdown in log processing while logstash instances were catching up on the backlog.

Note that the length of this incident was a contributing factor in the UDP loss, a shorter reoccurrence (20min) happened on Feb 11th due to repool of db1118 but resulted in no UDP loss on the logstash side as instances were able to catch up. A deeper understanding of logstash performance characteristics is needed as well.

Links to relevant documentation

- DB depooling: https://wikitech.wikimedia.org/wiki/MariaDB/troubleshooting

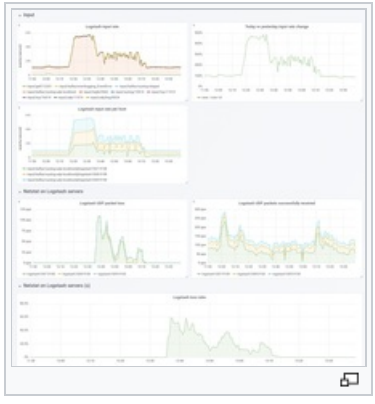

Actionables

- A MediaWiki dependency being down (single database host in this case) should not cause log spam/overload bug T215611
- The logging pipeline will need some additional spam / ratelimit protection bug T215900
- Better understanding of Logstash performance bug T215904

Category: Incident documentation