



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#).  
Please help us clean up older documentation referring to tools.wmflabs.org!

# Incident documentation/20190416-varnish

[< Incident documentation](#)

## Contents [\[hide\]](#)

- 1 Summary
  - 1.1 Impact
  - 1.2 Detection
- 2 Timeline
- 3 Conclusions
  - 3.1 What went well?
  - 3.2 What went poorly?
  - 3.3 Where did we get lucky?
- 4 Links to relevant documentation
- 5 Actionables

## Summary

For approximately an hour, the traffic layer served bursts of 503 errors: up to ~50k/minute for several minutes at a time. It is unclear why this happened, and whether misbehavior at the traffic layer or at the appserver layer was actually at fault.

## Impact

Approximately 553,000 HTTP 503 errors were served across all sites.

<https://logstash.wikimedia.org/goto/accfe83bffa587f460110942361af4a1>

## Detection

Automated monitoring (Icinga alerts on traffic availability) plus multiple staff/user reports in #wikimedia-operations.

## Timeline

All times in UTC.

- 18:24: HTTP 503 error rate begins to rise **OUTAGE BEGINS**
- 18:26: first alert from Icinga

```
<+icinga-wm> PROBLEM - HTTP availability for Varnish at esams on icingal001
is CRITICAL: job=varnish-text site=esams
https://grafana.wikimedia.org/dashboard/db/frontend-traffic?
panelId=3&fullscreen&refresh=1m&orgId=1
```

- 18:52: jynus phones bblack
- 18:55: cdanis depools cp1085 after looking at [Varnish mailbox lag console](#) and <https://phabricator.wikimedia.org/T145661>
- 19:04: bblack performs varnish-backend-restart on cp1085
- 19:07: bblack performs varnish-backend-restart on cp1083
- 19:08: 503s taper off to 0 **OUTAGE ENDS**

Possibly-relevant other graphs:

- Spikes of connections from Varnish backends to appservers correlate with the 503s: <https://grafana.wikimedia.org/d/000000439/varnish-backend-connections?orgId=1&from=1555438200000&to=1555442400000>
- Some very-long-running **WDQS** queries also correlate: <https://grafana.wikimedia.org/d/000000489/wikidata-query-service?panelId=26&fullscreen&orgId=1&from=1555438200000&to=1555442400000>

Main page  
Recent changes  
Server admin log (Prod)  
Server admin log (RelEng)  
Deployments  
SRE/Operations Help  
Incident status

Cloud VPS & Toolforge  
Cloud VPS documentation  
Toolforge documentation  
Request Cloud VPS project  
Server admin log (Cloud VPS)

Tools  
What links here  
Related changes  
Special pages  
Permanent link  
Page information  
Cite this page

Print/export  
Create a book  
Download as PDF  
Printable version

It is unclear what of the above are symptoms vs causes.

## Conclusions

---

### What went well?

- automated monitoring detected the incident

### What went poorly?

- unable to root-cause incident

### Where did we get lucky?

- Whatever was causing the issue stopped happening.
- The outage was not more widespread

## Links to relevant documentation

---

*Where is the documentation that someone responding to this alert should have (runbook, plus supporting docs). If that documentation does not exist, there should be an action item to create it.*

## Actionables

---

- Continue working on moving to ATS. Similar incidents have happened before, and continuing to investigate these failures of Varnish is not a good use of time.

Category: [Incident documentation](#)

This page was last edited on 19 April 2019, at 17:30.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.

[Privacy policy](#) [About](#)

[Disclaimers](#) [Code of Conduct](#) [Developers](#) [Statistics](#) [Cookie statement](#) [Mobile view](#)

[Wikitech](#)

