Main page
Recent changes
Server admin log (Prod)
Server admin log (RelEng)
Deployments
SRE/Operations Help
Incident status

Cloud VPS & Toolforge
Cloud VPS documentation
Toolforge documentation
Request Cloud VPS project
Server admin log (Cloud VPS)

Tools
What links here
Related changes
Special pages
Permanent link
Page information
Cite this page

Print/export
Create a book
Download as PDF
Printable version

Incident documentation/20190503-varnish

< Incident documentation

Summary

The traffic layer in eqiad reported various HTTP failures fetching from the application layer. FetchError seemed different from the usual varnish-be scalability issue (e.g. not the same as in Incident documentation/20190416-varnish). Instead of "Could not get storage" we were getting "HTC status -1" and (less frequently) "http format error". See https://phabricator.wikimedia.org/P8470 and https://phabricator.wikimedia.org/P8469 respectively. HTC status -1 happens when the function HTC_RxStuff returns HTC_S_EOF. The "HTTP format error" case was verified to *not* be happening at the appserver layer. This initially led us to think of a potentially different issue. However, the varnish backends throwing 503 errors were confirmed to consistently be those running for the longest amount of time, and restarting them solved the problem.

Impact

Approx 900,000 HTTP 503s served, mostly between 05:00 and 05:40.

Detection

Automated: Icinga paged.

Timeline

All times in UTC.

- 05:01: 503 errors begin at a slow but increasing rate **OUTAGE BEGINS**
- 05:09 Icinga pages about LVS in eqsin and codfw
- 05:32 Joe restarted varnish backend on cp1077
- 05:33 failed fetches move to cp1085
- 05:37 ema summoned
- 05:41 failed fetches from cp1085 disappear
- 05:51 failed fetches move to cp1089
- 06:00 failed fetches from cp1089 disappear **OUTAGE ENDS**
- 07:03 failed fetches from cp1089 return **OUTAGE BEGINS**
- 07:16 ema restarts varnish backend on cp1089
- 07:20 error rate returns to normal **OUTAGE ENDS**

Graphs:

- https://grafana.wikimedia.org/d/000000479/frontend-traffic?orgId=1&from=1556857800000&to=1556868300000
- https://logstash.wikimedia.org/goto/71cf7b8e8fde9e790d999276b7a57d3d

Conclusions

- The on-disk Varnish storage backend does not scale.

Actionables

- There is unfortunately no action that can be taken immediately. In the medium-term, our strategy is moving all cache backends to Apache Traffic Server.

Category: Incident documentation