



Toolforge webservices are in the final stages of [migrating to the toolforge.org domain](#) .
Please help us clean up older documentation referring to tools.wmflabs.org!

Incident documentation/20200604-cloud-private-repo

[< Incident documentation](#)**document status:** draft

Tracked in Phabricator
[Task T254491](#)

Summary

A change was deployed to puppet which inadvertently deleted the private repo from all puppet backend servers and puppet standalone servers. A few standalone servers in the Cloud environment maintain secrets by applying local commits to the labs/private repo. This event caused all secrets to be deleted required manual restoration

Impact: Any cloud environments which had added private secrets would have reverted to using the dummy secrets in the labs/private repo

Timeline

All timelines are on 2020-06-04 and are UTC

- 10:12: merge [change](#) to puppet-merge
- 10:12: **OUTAGE Begins** Once this change is merged the private repo will be removed the next time puppet is run (anytime between now and 30 mins)
- 10:32: [change is reverted](#) for unrelated reason, to fix a number of syntax errors
- 10:37: [change redeployed](#) with syntax errors fixed
- 10:58: jbond realises the private repo has been erroneously removed from stand-alone masters and [applies a fix](#)
- 10:58: jbond unaware real secrets where stored in some private repos did not realize the changes also affected puppet masters on WMCS projects leading to this incident
- 12:24: SAL !log did not work in #wikimedia-operations (worked at 12:18).
- 12:28: Arturo notices SAL does not work anymore (in #wikimedia-cloud)
- 12:32 <arturo> we don't have any [local] commit in labs/private in tools-puppetmaster-02
- 12:37: cloud engineers notice missing data in tools private repo and inquire about recent changes
- 12:42: confirmation that all private commits had been lost
- 12:55: explored option on using a temporary copy of the git repo created by git-sync , however that script deletes the temporary copy
- 12:58: start investigating if we can use /var/log/puppet.log to recover lost secrets
- 13:05: Ensure puppet is disabled on all cloud nodes
- 13:09: efforts made to use block level recovery to save data to an nfs mount
- 13:13: Bryan makes Antoine aware of the issue in #wikimedia-releng which would affect the CI and deployment-prep puppetmasters
- 13:18: Antoine backup private.git on integration/deployment-prep, disable puppet on them.
- 13:29: [confirmation](#) that deployment-prep and integrations where uneffected due to a merge conflict causing puppet updates to fail
- 13:39: Start meet up call to discuss next steps (Bryan, Arturo, John, Antoine [just at the beginning])
- 13:40: add toolforge-k8s-prometheus private key
- 13:43: reset root@wmflabs.org password for Project-proxy-dns-manager
- 13:45: start collection a copy of /var/log/puppet.log from all servers using cloud cumin

Contents [\[hide\]](#)

- [Summary](#)
- [Timeline](#)
- [Detection](#)
- [Conclusions](#)
 - [4.1 What went well?](#)
 - [4.2 What went poorly?](#)
 - [4.3 Where did we get lucky?](#)
 - [4.4 How many people were involved in the remediation?](#)
- [Links to relevant documentation](#)
- [Actionables](#)

[Main page](#)
[Recent changes](#)
[Server admin log \(Prod\)](#)
[Server admin log \(RelEng\)](#)
[Deployments](#)
[SRE/Operations Help](#)
[Incident status](#)

[Cloud VPS & Toolforge](#)

[Cloud VPS documentation](#)

[Toolforge documentation](#)

[Request Cloud VPS project](#)

[Server admin log \(Cloud VPS\)](#)

[Tools](#)

[What links here](#)

[Related changes](#)

[Special pages](#)

[Permanent link](#)

[Page information](#)

[Cite this page](#)

[Print/export](#)

[Create a book](#)

[Download as PDF](#)

[Printable version](#)

- 14:00: Start producing results from puppet.log files
- 14:03: commit Elasticsearch users and passwords to tools
- 14:19: Add keepalived password to tools
- 14:36: add k8s/kubeadm encryption key
- 14:52: add toolsview mysql password
- 15:07: Add docker private information to tools.
- 15:25: add puppet and puppetdb related secrets
- 15:29: use scp to copy all puppet.log files locally and confirm we have all secrets
- 16:14: add private password for tools-dns-manager for acme-chief
- 16:16: secrets for the acme-chief tools account
- 16:39: failry confident all 'urgent' breakages are now resolved
- 16:40 (Voila) **OUTAGE ENDS**

Detection

The issue was noticed by a member of the Cloud Services team

Copy the relevant alerts that fired in this section.

Did the appropriate alert(s) fire? Was the alert volume manageable? Did they point to the problem with as much accuracy as possible?

TODO: If human only, an actionable should probably be to "add alerting".

Conclusions

What went well?

- Cloud services and SRE foundations worked well to resolve the issue

What went poorly?

- The environment was broken for ~150 Minutes before being detected
- no backups
- lack of knowledge: engineer performing the original change was unaware the private repository was used in this manner and WMCS was not required to +1 in the code review

Where did we get lucky?

- deployment-prep and integrations puppet masters are automatically rebasing the puppet.git repository. Luckily there were merge conflicts on each of them that prevented the faulty change from being automatically pulled and deployed.
- we were able to restore the secrets from the `puppet.log` file

How many people were involved in the remediation?

- 4 SRE engineers troubleshooting the issue plus 1 incident commander

Links to relevant documentation

Add links to information that someone responding to this alert should have (runbook, plus supporting docs). If that documentation does not exist, add an action item to create it.

References

- IRC logs for the day:
 - <http://bots.wmflabs.org/logs/%23wikimedia-cloud/20200604.txt>
 - <http://bots.wmflabs.org/logs/%23wikimedia-operations/20200604.txt>
 - <http://bots.wmflabs.org/logs/%23wikimedia-releng/20200604.txt>

Actionables

Create a list of action items that will help prevent this from happening again as much as possible. Link to or create a Phabricator task for every step.

- To do #1 (TODO: Create task)

- [To do #2 \(TODO: Create task\)](#)

[TODO: Add the #Wikimedia-Incident-Prevention](#) [Phabricator](#) tag to these tasks.

Categories: [Incident documentation](#) | [Incident documentation drafts](#)

This page was last edited on 10 June 2020, at 16:02.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. See [Terms of Use](#) for details.

[Privacy policy](#) [About](#)

[Disclaimers](#) [Code of Conduct](#) [Developers](#) [Statistics](#) [Cookie statement](#) [Mobile view](#)

[Wikitech](#)

